

DIC Project Report

CSE 587B: Data Intensive Computing | Instructor: Dr. Shamsad Parvin

Name	UBIT	Person Number
Vivian Vincent Dmello	vivianvi	50540941
Chaitanya Deepak Yeole	cyeole	50535530
Vangala Eshwar Sai Prithvi Kiran	eshwarsa	50469012

Dataset: <https://www.kaggle.com/datasets/victorsoeiro/netflix-tv-shows-and-movies>

PHASE 1

1. Problem Statement:

Given a streaming platform such as Netflix, we want to identify which kinds of films go on to have great ratings. Based on things such as Run time, Genre, Actors, IMDB score and number of Votes, we want to see what we can use to predict movies with a good score.

Questions we are trying to answer:

- What determines this score? The Genre, Actors, Number of people who Voted/Watched it, the Run Time, etc?
- Which of these has the most value when it comes to determining this rating?
- What kind of impact does having different combinations of Genres have?

A. Why is this a significant problem?

There is an ocean of content out in the world when it comes to watchable material. Movies and TV Shows are just a few among the many ways people consume media. These attention capturing moving pictures aren't just there to entertain us but also teach us many things. However, the real world is one in which only the popular Shows and Movies come out on top. We cannot have fantastic media to watch without investors paying the bills and investors are drawn to profit. Also, the base desire of any consumer is to have not just one amazing show to consume, but to have multiple such shows readily available. Just from these points we can see that there is an intense need to be able to filter out the best possible stories and provide those to consumers as they are the most likely to be watched. After all many people pay attention to a Movies 'score' or 'rating' before they try to watch it!

B. What is the potential of your project to contribute to the problem domain and why is it crucial?

Our project can help contribute to the problem of 'selection' and 'prediction'. Our idea is to be able to use the readily available data to be able to choose from among sample movies the ones which would be the best rated and therefore most valued. Being able to choose this would involve looking at statistics of what makes a Movie good from the dataset we have chosen and using ML to be able to train a model to be able to pick out similar such shows. This would give a clearer idea on which films can be chosen and pushed to consumers.

2. Data Sources:

Our dataset is taken from Kaggle. This dataset is one which originally has been split into 2 different datasets – Titles and Credits Datasets. Titles Dataset has around 5k+ records with 15 features.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P
		title	type	description	release	age	runtime	genres	production	seasons	imdb_id	imdb_sc	imdb_votes	tmdb_pop	tmdb_sc	
1		Inception	MOVIE	Cobb, a skill	2010	PG-13	148	[action, 'sci-fi', 'music', 'thriller']	[US, 'GB']		tt1375666	8.8	2294231	108.284	8.4	
2		Forrest Gump	MOVIE	A man with	1994	PG-13	142	[drama, 'romance']	[US]		tt0109830	8.8	2021343	63.449	8.478	
3		The Dark Knight	MOVIE	Following the	2012	PG-13	x	[thriller, 'action', 'drama', 'crime']	[US]		tt1345836	8.4	1669067	91.76	7.768	
4		Se7en	MOVIE	Two homicide	1995	R	127	[crime, 'thriller', 'drama']	[US]		tt0114369	8.6	1606270	51.633	8.352	
5		Django Unchained	MOVIE	With the help	2012	R	165	[western, 'drama']	[US]		tt1853728	8.4	1472668	66.924	8.15	
6		The Departed	MOVIE	To take down	2006	R	151	[drama, 'thriller', 'crime', 'action']	[US]		tt0407887	8.5	1296244	33.795	8.2	
7		Titanic	MOVIE	101-year-old	1997	PG-13	194	[drama, 'romance']	[US]		tt0120338	7.9	1146825	155.683	7.878	
8		L.A. Confidential	MOVIE	L.A. Confidential	1994	R	111	[crime, 'drama', 'action', 'thriller', 'neo-noir']	[US, 'FR']		tt0110413	8.5	1137357	51.406	8.323	
9		GoodFellas	MOVIE	The true story	1990	R	145	[drama, 'crime']	[US]		tt0099685	8.7	1131681	50.387	8.463	
10		Catch Me If You Can	MOVIE	A true story	2002	PG-13	141	[drama, 'crime']	[US]		tt0264464	8.1	952602	72.321	8	
11		Snatch	MOVIE	Unscrupulous	2000	R	103	[crime, 'comedy']	[US, 'GB']		tt0208092	8.3	841435	25.203	7.8	
12		Taxi Driver	MOVIE	A mentally	1976	R	114	[drama, 'crime']	[US]		tt0075314	8.2	808582	40.965	8.179	
13		The Imitation Game	MOVIE	Based on the	2014	PG-13	113	[thriller, 'drama', 'war']	[GB, 'US']		tt2084970	8	756254	49.673	8	
14		I Am Legend	MOVIE	Robert Neville	2007	PG-13	101	[sci-fi, 'thriller', 'action', 'drama']	[US]		tt0480249	7.2	745093	69.139	7.192	
15		Full Metal Jacket	MOVIE	A pragmatic	1987	R	117	[war, 'drama']	[US, 'GB']		tt0093058	8.3	729274	32.937	8.1	
16		Casino Royale	MOVIE	Le Chiffre, a	2006		144	[action, 'thriller', 'european']	[DE, 'GB', 'US', 'BS', 'CZ']		tt0381061	8	649676	40.164	7.518	
17		The Amazing Spider-Man	MOVIE	Peter Parker	2012	PG-13	136	[action, 'fantasy', 'sci-fi']	[US]		tt0948470	6.9	646119	159.056	6.7	
18		Sherlock Holmes	MOVIE	Eccentric	2009	PG-13	129	[crime, 'thriller', 'action']	[AU, 'DE', 'GB', 'US']		tt0988045	7.6	625708	43.297	7.2	
19		Dunkirk	MOVIE	The story of	2017	PG-13	107	[drama, 'war', 'action', 'thriller', 'historical']	[FR, 'NL', 'GB', 'US']		tt5013758	7.8	169645	36.684	7.5	
20		Argo	MOVIE	As the Iranian	2012	R	120	[thriller, 'drama']	[GB, 'US']		tt1024648	7.7	600392	18.069	7.3	
21		The Hatefufu Movie	MOVIE	Bounty hunter	2015	R	188	[thriller, 'western', 'drama', 'crime']	[US]		tt3460252	7.8	570138	26.531	7.737	
22		Blade Runner	MOVIE	Thirty years	2017	R	164	[sci-fi, 'action', 'drama', 'thriller']	[CA, 'HU', 'MX', 'ES', 'GB', 'US']		tt1856101	8	549712	67.788	7.5	
23		Blood Diamond	MOVIE	An ex-military	2006	R	143	[drama, 'thriller', 'action']	[DE, 'US']		tt0450259	8	541547	47.39	7.5	
24		Monty Python and the Holy Grail	MOVIE	King Arthur	1975	PG	91	[fantasy, 'action', 'comedy']	[GB]		tt0071853	8.2	534486	15.461	7.811	
25		Nightcrawler	MOVIE	When Lou	2014	R	118	[thriller, 'crime', 'drama']	[US]		tt2872718	7.8	531779	57.276	7.7	
26		It	MOVIE	In a small town	2017	R	135	[horror, 'fantasy', 'thriller']	[US]		tt1396484	7.3	527626	95.827	7.234	
27		Troy	MOVIE	In year 1250	2004	R	163	[action, 'history', 'war', 'drama', 'euro']	[MT, 'GB', 'US']		tt0332452	7.3	527447	48.407	7.1	

Credits data set on the other hand has over 70k+ records with 5 features.

person_id	id	name	character	role
3748	tm84618	Robert De Niro	Travis Bickle	ACTOR
14658	tm84618	Jodie Foster	Iris Steensma	ACTOR
7064	tm84618	Albert Brooks	Tom	ACTOR
3739	tm84618	Harvey Keitel	Matthew 'Sport' Higgins	ACTOR
48933	tm84618	Cybill Shepherd	Betsy	ACTOR
32267	tm84618	Peter Boyle	Wizard	ACTOR
519612	tm84618	Leonard Nimoy	Ch. ACTOR	
29068	tm84618	Dianne Wiest	Concession	ACTOR
519613	tm84618	Gino Arditi	Policeman	ACTOR
3308	tm84618	Martin Scorsese	Passenger	ACTOR
43791	tm84618	Murray Close	Time	ACTOR
519614	tm84618	Richard Hill	Secret Serv	ACTOR
519615	tm84618	Bill Minkin	Tom's Assis	ACTOR
82426	tm84618	Bob Marlin	Madison	ACTOR
20935	tm84618	Victor Argo	Mello, Deli	ACTOR
7753	tm84618	Joe Spinell	Personell	ACTOR
43279	tm84618	Robinson F. Anger	Blad	ACTOR
519618	tm84618	Brenda Dick	Soap Oper	ACTOR
8424	tm84618	Norman M. Charlie	T	ACTOR
20447	tm84618	Harry North	Doughboy	ACTOR
519618	tm84618	Harlan Carr	Campaign	ACTOR
49567	tm84618	Steven Price	Andy - Gun	ACTOR
15551	tm84618	Peter Savoy	The John	ACTOR
475303	tm84618	Nicholas Sp	Palantine's	ACTOR
43540	tm84618	Ralph S. St	T.V. Interv	ACTOR
466061	tm84618	Annie Gage	Campaign	ACTOR
212023	tm84618	Carson Gra	Political	ACTOR

Link for the datasets - <https://www.kaggle.com/datasets/victorsoeiro/netflix-tv-shows-and-movies>

3. Data Cleaning/Processing:

1. Merged the titles and credits dataset:

In the initial step, we have merged the two datasets we have taken on the basis of the ID field found in both datasets. First we have read the datasets and stored into dataframes as below:

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

[2] credits = (pd.read_csv(r"credits.csv"))
credits

    person_id    id      name  character  role
0         3748  tm84618  Robert De Niro  Travis Bickle  ACTOR
1        14658  tm84618    Jodie Foster    Iris Steensma  ACTOR
2         7064  tm84618  Albert Brooks         Tom  ACTOR
3         3739  tm84618  Harvey Keitel  Matthew 'Sport' Higgins  ACTOR
4        48933  tm84618  Cybill Shepherd         Betsy  ACTOR
...         ...      ...
77796    736339  tm1059008  Adelaïda Buscato  Maria Paz  ACTOR
77797    399499  tm1059008  Luz Stela Luengas  Karen Bayona  ACTOR
77798    373198  tm1059008    Inés Prieto         Fanny  ACTOR
77799    378132  tm1059008  Isabel Gaxiola         Cécilia  ACTOR
77800   1950416  tm1059008  Julian Gaviria         NaN  DIRECTOR
77801 rows x 5 columns
  
```

```

[4] movies = (pd.read_csv(r"titles.csv"))
movies

    id  title  type  description  release_year  age_certification  runtime  genres  production_countries  seasons  ind
0   NaN  Five Came Back: The Reference Films  SHOW  This collection includes 12 World War II-era p...  1945.0  TV-MA  51  [documentary]  [US]  1.0
1  tm84618  Taxi Driver  MOVIE  A mentally unstable Vietnam War veteran works ...  NaN  R  114  [drama, 'crime']  [US]  NaN  8007
2  tm154986  Deliverance  MOVIE  Intent on seeing the Cahulawassee River before...  NaN  R  109  [drama, 'action', 'thriller', 'adventure']  [US]  NaN  8006
3  tm127384  Monty Python and the Holy Grail  MOVIE  King Arthur, accompanied by his squire, recu...  NaN  PG  91  ['fantasy', 'action', 'comedy']  [GB]  NaN  8007
4  tm120801  The Dirty Dozen  MOVIE  12 American military prisoners in World War II...  NaN  NaN  150  [war, 'action']  [GB, US]  NaN  8006
...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...
  
```

Next we have stored the Actor names as lists within a dictionary keeping the ID for all those actors as the key value:

The screenshot displays the dictlymb IDE interface. At the top, the title bar reads "dictlymb" with a star icon. Below it is a menu bar with options: "File", "Edit", "View", "Insert", "Runtime Tools", and "Help". A status bar on the right shows icons for "Comment", "Share", and a red circle with a white "X".

The left sidebar contains a "Files" panel with a tree view showing a project structure:

- sample_data
 - credits.csv
 - titles.csv

The main editor area is titled "+ Code + Text" and contains a Python script:

```

names_dict = {}
for id in credits['id'].unique():
    # Use Pandas to filter and extract names for each person_id
    names_list = credits[credits['id'] == id]['name'].tolist()

    # Store the names list in the dictionary with person_id as the key
    names_dict[id] = names_list

names_dict

```

The script defines a dictionary named `names_dict` by iterating over the unique values of the `id` column in the `credits` DataFrame. For each `id`, it filters the DataFrame to get the corresponding names and stores them in the dictionary. The final state of `names_dict` is displayed below the code:

```

{'11558': ['Nissar Khan',
          'Raghu Godse',
          'Rajesh Khanna',
          'Rajiv Gupta',
          'Tigmanshu Dhulia'],
 '1215846': ['Eraman Mishra',
            'Humaine Malik',
            'Dhruv Raut',
            'Kay Kay Menon',
            'Deepak Tijori',
            'Muhammad Zareesh Ayyub',
            'Kunal Deshmukh'],
 '122022': ['Lucie Soderstrand',
            'Sophie Robinson',
            'Jan Soderstrand',
            'Hente Soderstrand',
            'David Lynch',
            'Sophie Robinson',
            'Lotje Soderstrand'],
 '1231218': ['Jim Gaffigan', 'Jay Chapman'],
 '1246607': ['Jeff Garmy', 'Rob Dipple'],
 '1257681': ['Khaleel Abul Naga',
            'Rajesh Khanna',
            'Rajiv Gupta',
            'Tigmanshu Dhulia',
            'Humaine Malik',
            'Dhruv Raut',
            'Kay Kay Menon',
            'Deepak Tijori',
            'Muhammad Zareesh Ayyub',
            'Kunal Deshmukh']}

```

At the bottom of the interface, a status bar indicates "Disk" usage and "81.26 GB available".

Now we have created a Dataframe containing the ID and list of Actor names:

The screenshot shows the dclipyb application interface. At the top, there is a menu bar with options: File, Edit, View, Insert, Runtime, Tools, and Help. Below the menu bar, there is a sidebar on the left with a file explorer showing a directory structure with 'sample_data', 'credits.csv', and 'titles.csv'. The main editor area displays a Jupyter Notebook cell with the following code:

```
combined_df = pd.DataFrame({'id': list(names_dict.keys()), 'actors': list(names_dict.values())})

# "combined_df" will contain the movie id and associated list of actors
combined_df
```

The output of the code is a DataFrame with two columns: 'id' and 'actors'. The DataFrame contains 5489 rows of movie data. The first few rows are:

	id	actors
0	tm84618	[Robert De Niro, Jodie Foster, Albert Brooks, ...]
1	tm154986	[Jon Voight, Burt Reynolds, Ned Beatty, Romyr...
2	tm127384	[Graham Chapman, John Cleese, Eric Idle, Terry...
3	tm120801	[Lee Marvin, Ernest Borgnine, Charles Bronson, ...]
4	ts22164	[Graham Chapman, Michael Palin, Terry Jones, E...

The DataFrame is displayed with 5489 rows and 2 columns. The status bar at the bottom indicates '11.25 GB available'.

Finally we have merged both datasets on the basis ID field as below:

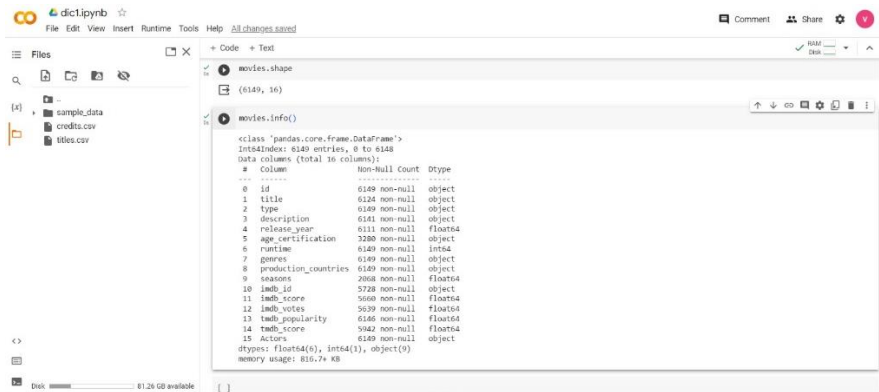
The screenshot shows the dclplaynb interface with a Jupyter Notebook. The code cell contains the following Python code:

```
merged_df = pd.merge(movies, 'combined_df', on='id', how='inner')
movies = merged_df
```

The output is a DataFrame with the following columns: id, title, type, description, release_year, age_certification, runtime, genres, production_countries, seasons, and imdb_id. The DataFrame contains 4 rows of movie data:

	id	title	type	description	release_year	age_certification	runtime	genres	production_countries	seasons	imdb_id
0	tm84618	Taxi Driver	MOVIE	A mentally unstable Vietnam War veteran works...		NaN	114	[drama, 'crime']	[US]	NaN	tt0075314
1	tm154086	Deliverance	MOVIE	Intent on seeing the Cahowawee River before...		NaN	100	[drama, 'action', 'thriller', 'europcan']	[US]	NaN	tt0068473
2	tm127384	Monty Python and the Holy Grail	MOVIE	King Arthur, accompanied by his squire, repul...		NaN	91	[fantasy, 'action', 'comedy']	[GB]	NaN	tt0071853
3	tm120801	The Dirty Dozen	MOVIE	12 American military prisoners in World War II...		NaN	150	[war, 'action']	[GB, US]	NaN	tt0061578

Below is the information of the final combined Dataframe:

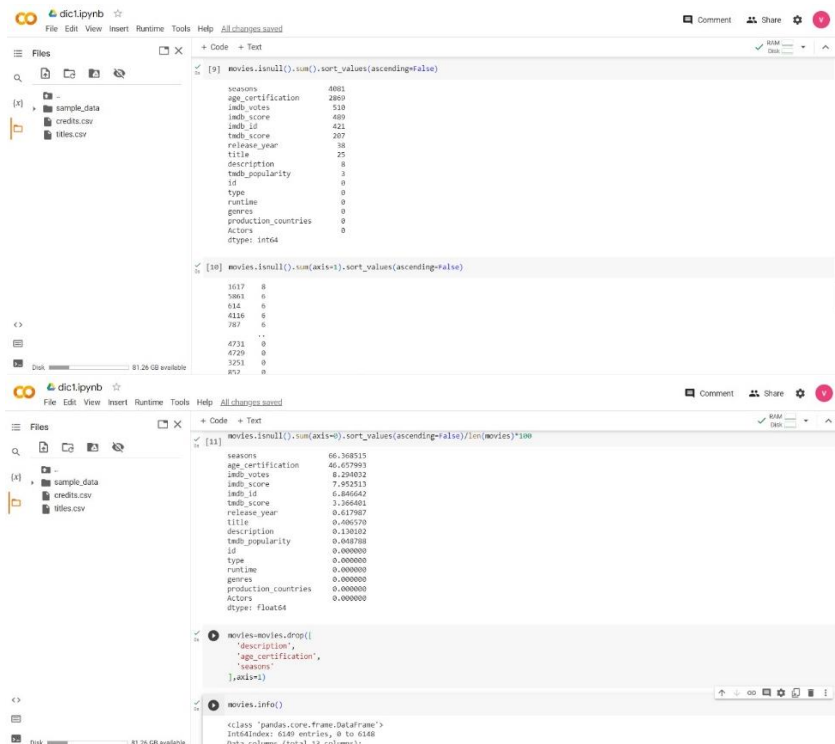


```
movies.shape
(6149, 16)

movies.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6149 entries, 0 to 6148
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   id                    6149 non-null  object  
 1   title                 6124 non-null  object  
 2   type                  6149 non-null  object  
 3   description            6141 non-null  object  
 4   release_year          6111 non-null  float64  
 5   age_certification      3300 non-null  object  
 6   runtime               6149 non-null  int64  
 7   genres                 6149 non-null  object  
 8   production_countries  6149 non-null  object  
 9   seasons               2068 non-null  float64  
10  imdb_id               5728 non-null  object  
11  imdb_score            5568 non-null  float64  
12  imdb_votes            5639 non-null  float64  
13  tmdb_popularity        6146 non-null  float64  
14  tmdb_score            5942 non-null  float64  
15  Actors                 6149 non-null  object  
dtypes: float64(6), int64(1), object(9)
memory usage: 816.7+ KB
```

2. Removed Description, age_certification and season columns:

For this step we have checked which columns are having the most count of NULL values as below:



```
[9] movies.isnull().sum().sort_values(ascending=False)
seasons          4081
age_certification 2869
imdb_votes        510
imdb_id           489
tmdb_id           421
tmdb_score        207
release_year       18
title              25
description         8
tmdb_popularity    3
id                 0
type               0
runtime            0
genres             0
production_countries 0
Actors             0
dtype: int64

[10] movies.isnull().sum(axis=1).sort_values(ascending=False)
1617  8
5801  6
614  6
4116  6
787  6
..
4731  0
4729  0
3251  0
812  0
dtype: int64

[11] movies.isnull().sum(axis=0).sort_values(ascending=False)/len(movies)*100
seasons          66.368515
age_certification 46.677993
imdb_votes        8.294032
imdb_id           7.952513
tmdb_id           6.846642
tmdb_score        3.366401
release_year      0.417907
title             0.406570
description       0.130102
tmdb_popularity   0.048788
id                0.000000
type              0.000000
runtime           0.000000
genres            0.000000
production_countries 0.000000
Actors            0.000000
dtype: float64

movies = movies.drop([
    'description',
    'age_certification',
    'seasons'
], axis=1)

movies.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6149 entries, 0 to 6148
Data columns (total 13 columns):
```

Based on the above statistics, we have removed the Description, Age_certification and Season columns as they are having high amount of NULL values and will also not be of much use while training the ML models.

```

movies = movies.drop(
    ['description',
     'age_certification',
     'seasons'],
    axis=1)

movies.info()

```

class 'pandas.core.frame.DataFrame'>
Int64Index: 6160 entries, 0 to 6158
Data columns (total 13 columns):
column non-null count dtype
0 id 6160 non-null object
1 title 6124 non-null object
2 type 6149 non-null object
3 release_year 6111 non-null float64
4 runtime 6149 non-null float64
5 genres 6149 non-null object
6 production_countries 6149 non-null object
7 imdb_id 5728 non-null object
8 imdb_score 5669 non-null float64
9 imdb_votes 5639 non-null float64
10 tmdb_popularity 6146 non-null float64
11 tmdb_score 5942 non-null float64
12 actors 6149 non-null object
dtypes: float64(3), int64(1), object(9)
memory usage: 672.5+ KB

3. Calculated percentage of empty rows, then removed empty rows (blank spaces or Nulls):

Here we are checking how many of the remaining columns have NULL values and using those columns, we have removed all the NULL rows as we require these columns later on.

```

round(movies.isnull().sum(), sort_values(ascending=False)/len(movies)*100,2)

imdb_votes      8.29
imdb_score       7.95
imdb_id         6.85
tmdb_votes       3.37
tmdb_score       0.62
release_year     0.41
title            0.05
tmdb_popularity  0.00
id              0.00
type            0.00
runtime         0.00
genres          0.00
production_countries 0.00
actors          0.00
dtype: float64

movies = movies[movies['imdb_votes'].notnull()]
movies = movies[movies['imdb_score'].notnull()]
movies = movies[movies['tmdb_id'].notnull()]
movies = movies[movies['tmdb_score'].notnull()]

round(movies.isnull().sum(), sort_values(ascending=False)/len(movies)*100,2)

release_year     0.62
title            0.25
id              0.00
type            0.00
runtime         0.00
genres          0.00

```

4. Filled in Missing values of Year column:

In this step we have checked Year column and replaced all its NULL values using the mean value for the Year column as we cannot have this field be empty:

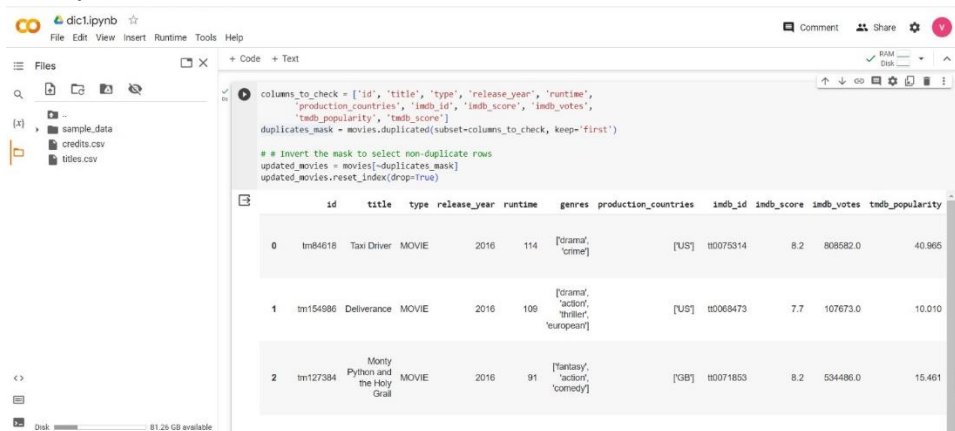
```

mean_year = movies['release_year'].mean()
movies['release_year'] = movies['release_year'].fillna(mean_year).astype(int)
movies

```

	id	title	type	release_year	runtime	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity
0	tm84618	Taxi Driver	MOVIE	2016	114	[drama, 'crime']	[US]	tt0075314	8.2	808582.0	40.965
1	tm154898	Deliverance	MOVIE	2016	109	[drama, 'action', 'thriller', 'adventure']	[US]	tt0068473	7.7	107673.0	10.010
2	tm127384	Monty Python and the Holy Grail	MOVIE	2016	91	[fantasy, 'action', 'comedy']	[GB]	tt0071853	8.2	534486.0	15.461
3	tm120801	The Dirty Dozen	MOVIE	2016	150	[war, 'action']	[GB, US]	tt0061578	7.7	72662.0	20.368

5. Removed duplicate rows:

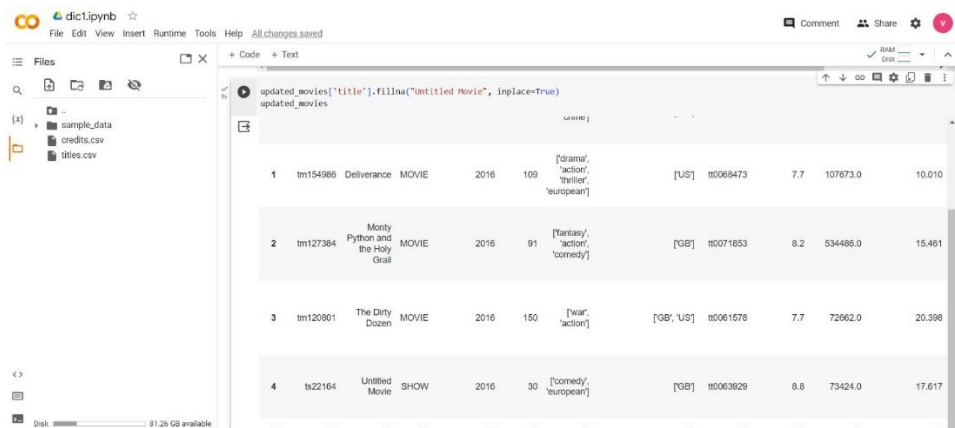


```
columns_to_check = ['id', 'title', 'type', 'release_year', 'runtime',  
                    'production_countries', 'imdb_id', 'imdb_score', 'imdb_votes',  
                    'tmdb_popularity', 'tmdb_score']  
duplicates_mask = movies.duplicated(subset=columns_to_check, keep='first')  
  
# Invert the mask to select non-duplicate rows  
updated_movies = movies[~duplicates_mask]  
updated_movies.reset_index(drop=True)
```

	id	title	type	release_year	runtime	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity
0	tm84618	Taxi Driver	MOVIE	2016	114	[drama, 'crime']	[US]	tt0075314	8.2	808582.0	40.965
1	tm154986	Deliverance	MOVIE	2016	109	[drama, 'action', 'thriller', 'european']	[US]	tt0068473	7.7	107673.0	10.010
2	tm127384	Monty Python and the Holy Grail	MOVIE	2016	91	[fantasy, 'action', 'comedy']	[GB]	tt0071853	8.2	534486.0	15.461

6. Movies which have Title field as empty, we have replaced as 'Untitled Movie':

For this step we have replaced the Movies with empty titles to have a common 'Untitled Movie' name. This we will keep as it is as we are more concerned about predicting the final rating. The movie name only serves as a way to differentiate different proposals as per the problem statement we have given.

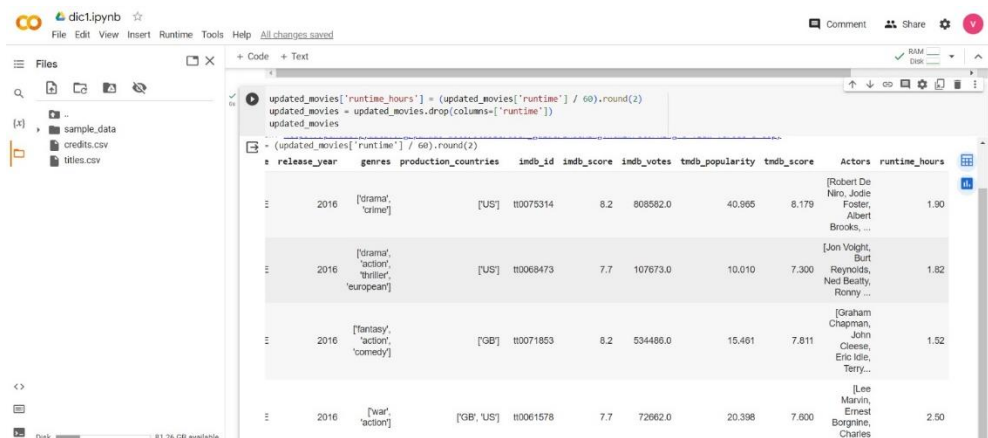


```
updated_movies['title'].fillna("Untitled Movie", inplace=True)
```

	id	title	type	release_year	runtime	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity
1	tm154986	Deliverance	MOVIE	2016	109	[drama, 'action', 'thriller', 'european']	[US]	tt0068473	7.7	107673.0	10.010
2	tm127384	Monty Python and the Holy Grail	MOVIE	2016	91	[fantasy, 'action', 'comedy']	[GB]	tt0071853	8.2	534486.0	15.461
3	tm120801	The Dirty Dozen	MOVIE	2016	150	[war, 'action']	[GB, US]	tt0061578	7.7	72662.0	20.398
4	tt22164	Untitled Movie	SHOW	2016	30	[comedy, 'european']	[GB]	tt0063629	8.8	73424.0	17.617

7. Converting run time from minutes to hours:

We have converted runtime from minutes to hours to allow for better feature scaling:

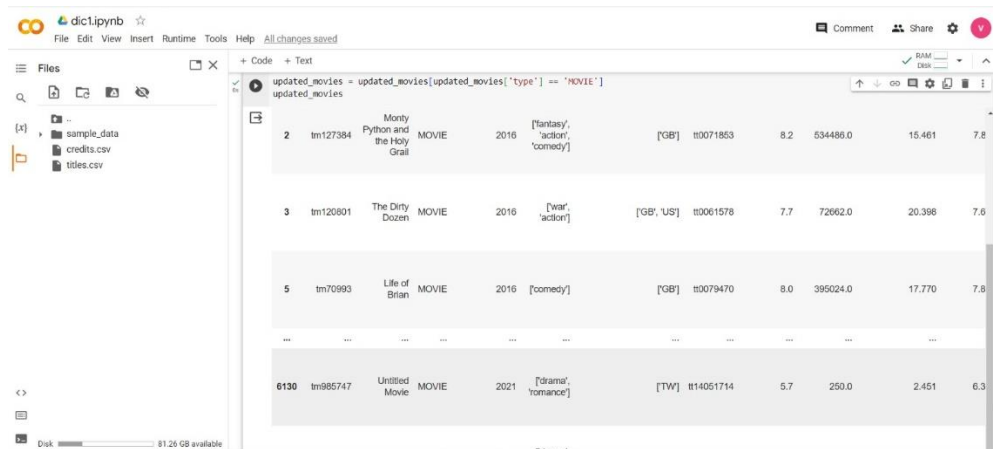


```
updated_movies['runtime_hours'] = (updated_movies['runtime'] / 60).round(2)  
updated_movies = updated_movies.drop(columns=['runtime'])  
updated_movies
```

```
(updated_movies['runtime'] / 60).round(2)
```

	release_year	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity	tmdb_score	Actors	runtime_hours
0	2016	[drama, 'crime']	[US]	tt0075314	8.2	808582.0	40.965	8.179	[Robert De Niro, Jodie Foster, Albert Brooks, ...]	1.90
1	2016	[drama, 'action', 'thriller', 'european']	[US]	tt0068473	7.7	107673.0	10.010	7.300	[Jon Voight, Burt Reynolds, Ned Beatty, Ronny ...]	1.82
2	2016	[fantasy, 'action', 'comedy']	[GB]	tt0071853	8.2	534486.0	15.461	7.811	[Graham Chapman, John Cleese, Eric Idle, Terry ...]	1.52
3	2016	[war, 'action']	[GB, US]	tt0061578	7.7	72662.0	20.398	7.800	[Lee Marvin, Ernest Borgnine, Charles ...]	2.50

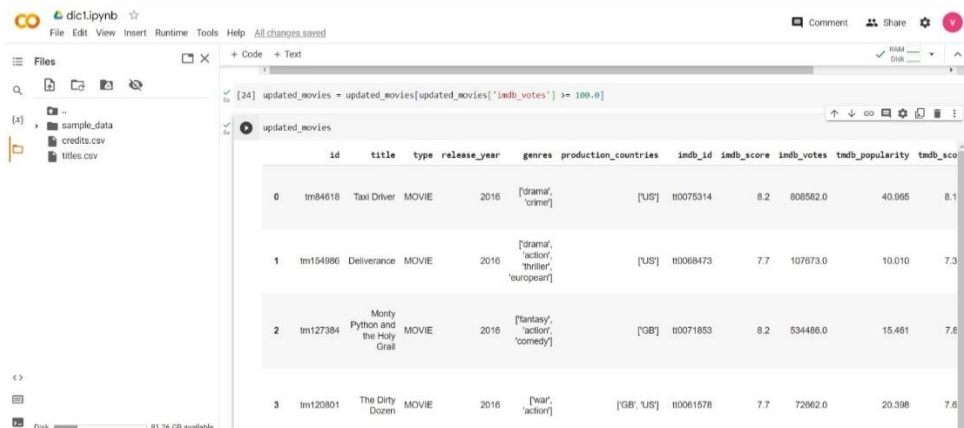
8. We are removing all the data regarding shows and keeping only movies:
For this step we are taking into consideration only Movies and rejecting TV shows to allow for better training in one specific field. Conflating both would lead to incorrect results as runtimes and actors would vary by a lot.



updated_movies = updated_movies[updated_movies['type'] == 'MOVIE']

	id	title	type	release_year	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity	tmdb_score
2	tm127384	Monty Python and the Holy Grail	MOVIE	2016	['fantasy', 'action', 'comedy']	['GB']	tt0071863	8.2	534486.0	15.461	7.8
3	tm120801	The Dirty Dozen	MOVIE	2016	['war', 'action']	['GB', 'US']	tt0061578	7.7	72662.0	20.398	7.6
5	tm70993	Life of Brian	MOVIE	2016	['comedy']	['GB']	tt0079470	8.0	395024.0	17.770	7.8
6130	tm985747	Untitled Movie	MOVIE	2021	['drama', 'romance']	['TW']	tt14061714	5.7	250.0	2.451	6.3

9. Removing outliers for Votes field:
Here we have removed outliers lesser than 100 votes from the IMDB votes field.

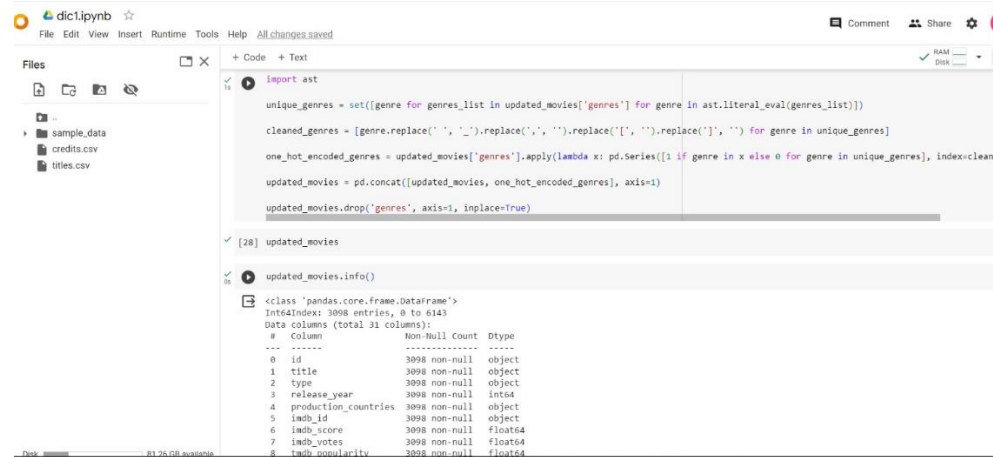


updated_movies = updated_movies[updated_movies['imdb_votes'] >= 100.0]

	id	title	type	release_year	genres	production_countries	imdb_id	imdb_score	imdb_votes	tmdb_popularity	tmdb_score
0	tm84618	Taxi Driver	MOVIE	2016	['drama', 'crime']	['US']	tt0075314	8.2	806562.0	40.865	8.1
1	tm154986	Deliverance	MOVIE	2016	['drama', 'action', 'thriller', 'horror']	['US']	tt0068473	7.7	107673.0	10.010	7.3
2	tm127384	Monty Python and the Holy Grail	MOVIE	2016	['fantasy', 'action', 'comedy']	['GB']	tt0071863	8.2	534486.0	15.461	7.8
3	tm120801	The Dirty Dozen	MOVIE	2016	['war', 'action']	['GB', 'US']	tt0061578	7.7	72662.0	20.398	7.6

10. Splitting genres through one-hot encoding and keeping only most relevant entries, ie with most frequency:

Here we have performed one-hot encoding on the Genre column to create a numerical value for each class of this categorical variable to allow it to be included for processing within the ML models:



```
import ast

unique_genres = set([genre for genres_list in updated_movies['genres'] for genre in ast.literal_eval(genres_list)])

cleaned_genres = [genre.replace(' ', '').replace(',','').replace('[','').replace(']', '') for genre in unique_genres]

one_hot_encoded_genres = updated_movies['genres'].apply(lambda x: pd.Series([1 if genre in x else 0 for genre in unique_genres], index=cleaned_genres))

updated_movies = pd.concat([updated_movies, one_hot_encoded_genres], axis=1)

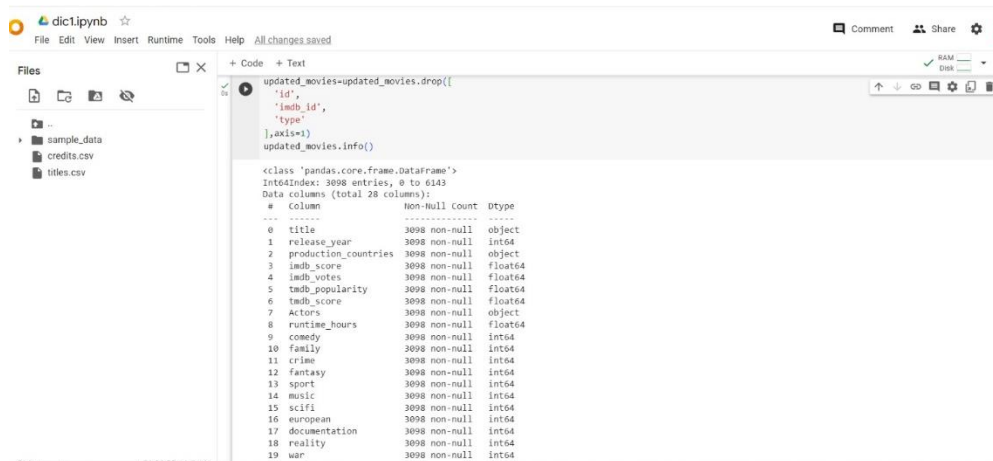
updated_movies.drop('genres', axis=1, inplace=True)
```

```
[28] updated_movies
```

```
updated_movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3098 entries, 0 to 6143
Data columns (total 31 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   id                  3098 non-null   object
 1   title               3098 non-null   object
 2   type                3098 non-null   object
 3   release_year        3098 non-null   int64
 4   production_countries 3098 non-null   object
 5   imdb_id             3098 non-null   object
 6   imdb_score          3098 non-null   float64
 7   imdb_votes          3098 non-null   float64
 8   tmdb_popularity     3098 non-null   float64
```

As the final step, we have also removed the columns which we deemed unnecessary such as ID, IMDB_ID and TYPE:



```
updated_movies=updated_movies.drop(['id', 'imdb_id', 'type'], axis=1)
updated_movies.info()
```

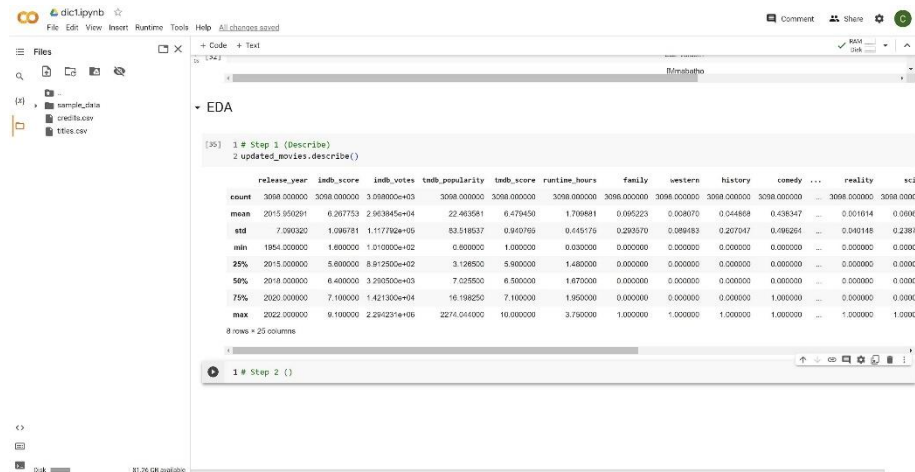
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3098 entries, 0 to 6143
Data columns (total 28 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   title               3098 non-null   object
 1   release_year        3098 non-null   int64
 2   production_countries 3098 non-null   object
 3   imdb_score          3098 non-null   float64
 4   imdb_votes          3098 non-null   float64
 5   tmdb_popularity     3098 non-null   float64
 6   tmdb_score          3098 non-null   float64
 7   Actors              3098 non-null   object
 8   runtime_hours       3098 non-null   float64
 9   comedy              3098 non-null   int64
10   family              3098 non-null   int64
11   crime               3098 non-null   int64
12   fantasy              3098 non-null   int64
13   sport               3098 non-null   int64
14   music               3098 non-null   int64
15   sci-fi              3098 non-null   int64
16   european            3098 non-null   int64
17   documentation       3098 non-null   int64
18   reality             3098 non-null   int64
19   war                 3098 non-null   int64
```

4. Exploratory Data Analysis:

As part of the EDA, we are performing the below steps:

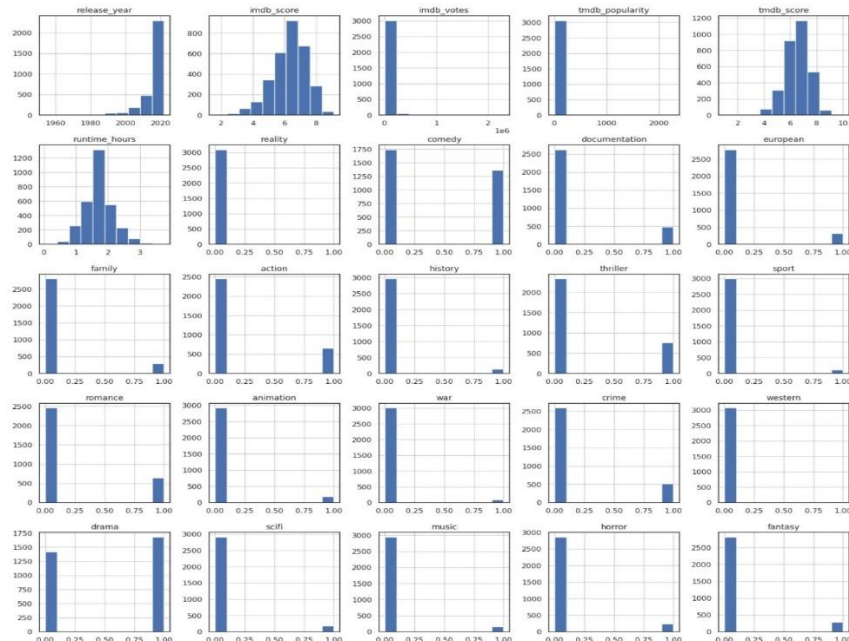
1. Checking measures of Central Tendency and Variance:

Using Describe() we are checking the Mean, Median, Standard Deviation, Maximum and Minimum value over the dataset to better understand its distribution:




2. Plotting Histogram:

Here we have plotted histograms to understand the distribution of all the fields. We can see that both score fields have normal distributions along with run-time hour fields. The genre fields are all one hot encoded to have only 1 value so for models which require prediction based on normalized values such as logistic regression, we will assign weighted value to each genre based on total frequency and combine all the separate fields into one to allow for better prediction for such type of models. For models requiring classification, we can use the remaining data and fields with good correlation to predict the genre of the particular movie.



3. Plotting Boxplots:


We have used boxplots to identify outliers in all the fields and as expected, there will not be anything in the genre fields however it is showing we have outliers in the votes field. We can remove these particular outliers going forward for a few specific models that we will utilize depending on the requirement however for regression analysis, it might be better to keep the values since the number of votes being greater signifies a greater weightage toward the film score being more solid.



From the above, we can see that the scores also have a few outliers although they should not pose much of a problem as there are not too many.

4. Plotting Correlation Matrix:

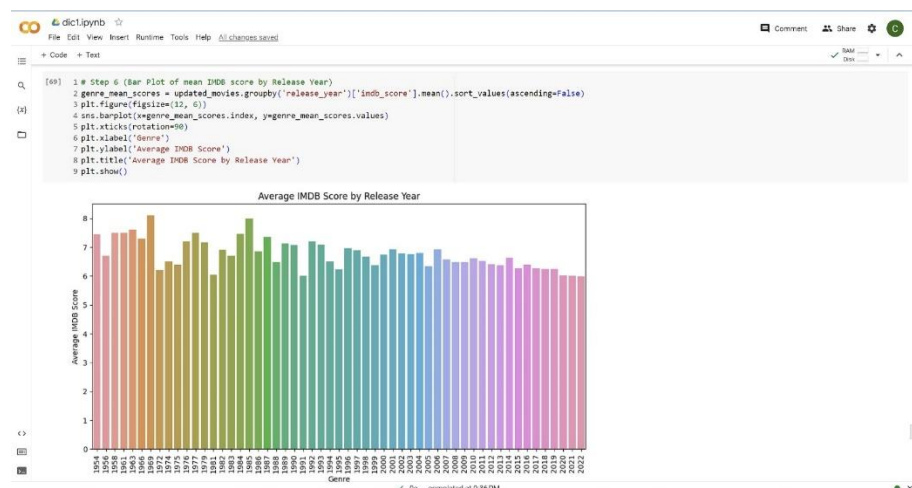
We have plotted a correlation matrix to check correlation between features. We can see that the IMDB score has good correlation with the TMDB score and a pretty good correlation with the IMDB votes. The other genre fields which were one hot encoded have differing levels of correlation with all the fields although there are particular genres which have good correlation with other genres.



Based on the relation between IMDB score and TMDB score we can see that we can use one system to more or less predict the score in the other system, if model is trained accurately.

5. Plotting bar plot:

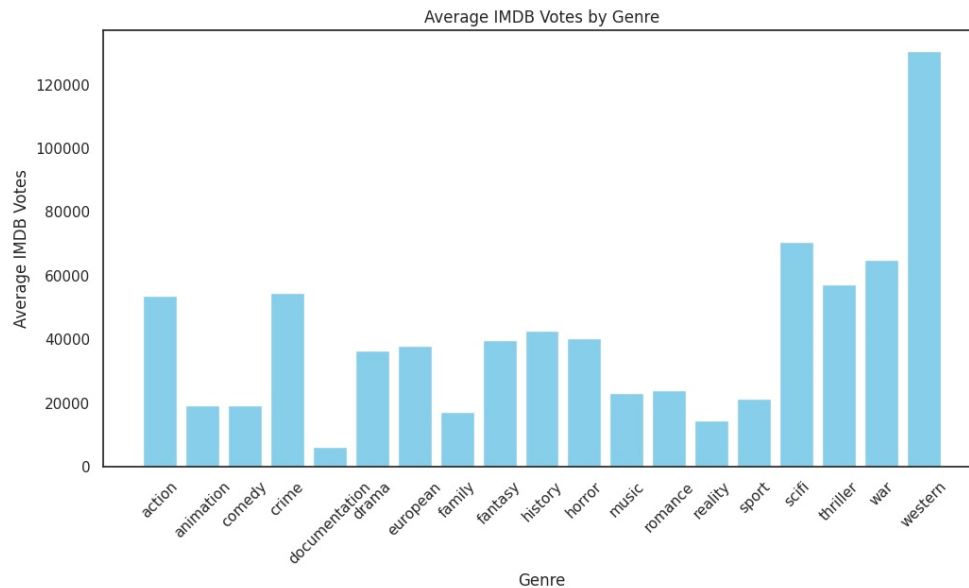
Here we have plotted a bar plot between IMDB score and year. We can see that the variation is quite little majority of scores being between 6 and 8 throughout the years.



6. Frequency Bar plot between each Genre and Votes:

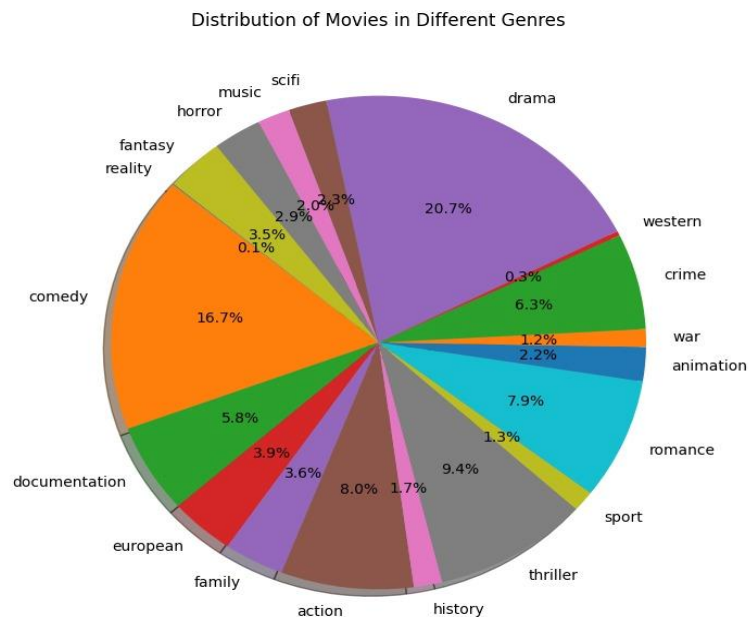
Here we are checking the total number votes per genre to understand which genres are the most common and therefore most popular. From below we can see that westerns are extremely popular followed by Scifi and War themed movies.

Based on this we should be able to identify the more popular genres through the number of votes.



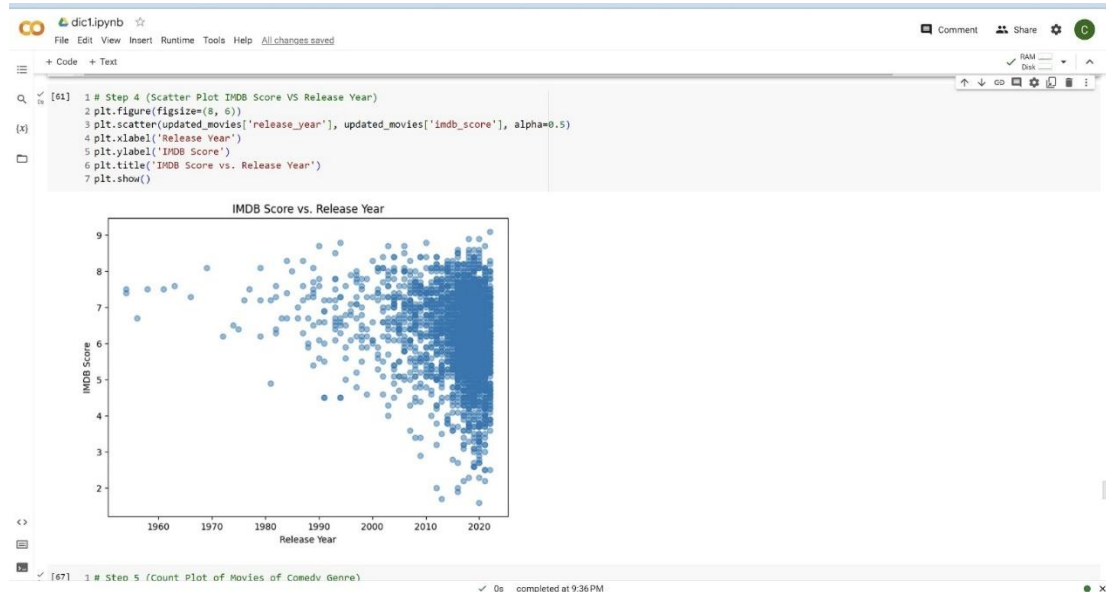
7. Pie chart to show percentage of all genres:

We can see the percentage of Genres for all of the films as below. From below we are able to see that drama and comedy are the most common followed by romance and action genres. Through this we can see that these are important Genres although further analysis is required.



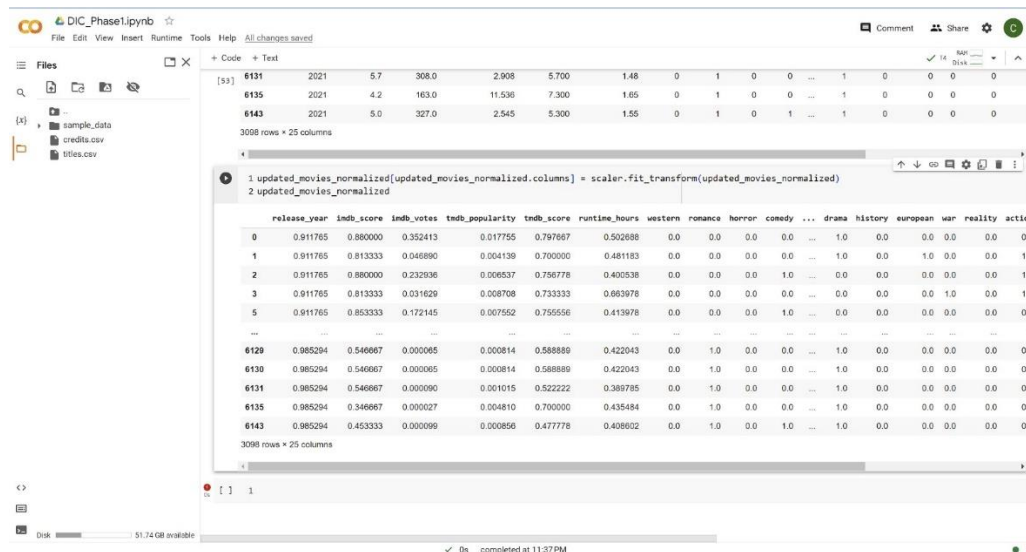
8. Scatter plot:

Here we are plotting a scatter plot between score and release year to check score throughout the years. We see that a huge amount of content is generated progressively throughout the years and the vast majority of content is concentrated in the 2000 to 2020 years. Based on this we will aim to remove the films which are lesser than 1980 to remove the outliers from that time and keep the rest.



9. Normalize the data:

Here we have normalized the data to allow for features such as Year and Population to be scaled properly and not cause overfitting due to their extreme values. Also, through this normalization, it will be easier to fit the training model for a few ML algorithms.



For this step we will create a separate column representing Genre as 1 value numerically. Here we wanted to create a column that is the sum of the genres with each genre being multiplied with a proportional weighted sum. It is our belief that this would be useful for at least one ML model in phase 2.

PHASE 2

Decision Tree Classifier:

Justification:

- Because of its simplicity of interpretation and capacity to manage intricate, non-linear relationships within the data, the Decision Tree method was selected.
- Movies can be good for a lot of varied reasons (genres, directors, budgets, etc.), and Decision Trees' non-parametric structure makes it possible to model these interactions well.

Work Done:

- A grid search was performed across multiple hyperparameters, such as tree depth and the minimal number of samples needed to split an internal node, to improve the Decision Tree.
- This was essential to maintaining the model's generalizability and avoiding overfitting.

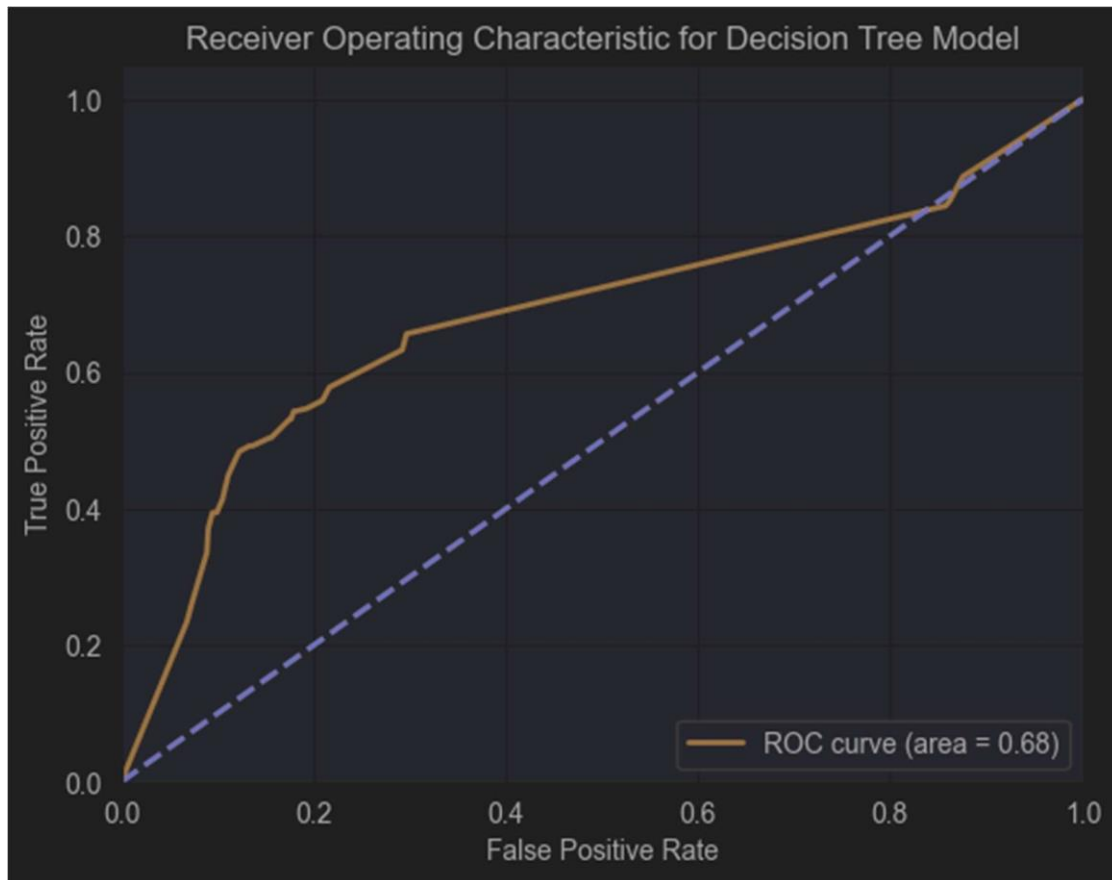
Model Effectiveness:

- After the model was evaluated, the Decision Tree obtained an AUC value of 0.6797 and an accuracy of roughly 76.24%.

```
Decision Tree Model Accuracy: 0.7623655913978494  
AUC Score: 0.6797077846810089
```


ROC Curve Analysis:

- With an AUC of 0.68, the Decision Tree model's ROC curve indicates a decent ability to distinguish between the classes. Given that the curve is closer to the diagonal line of no discrimination than the ideal top-left corner, this shows that there is some opportunity for improvement.



Logistic Regression:

Justification:

- Because of its effectiveness in predicting probabilities—a critical skill when dealing with binary outcomes like good vs. bad movies—Logistic Regression was chosen. It can be used as a benchmark for more complicated models, providing a linear decision boundary as a baseline.

Work Done:

- To prevent both overfitting and underfitting in Logistic Regression, the regularization strength (C) was adjusted to strike the ideal balance between training data fit and model complexity.

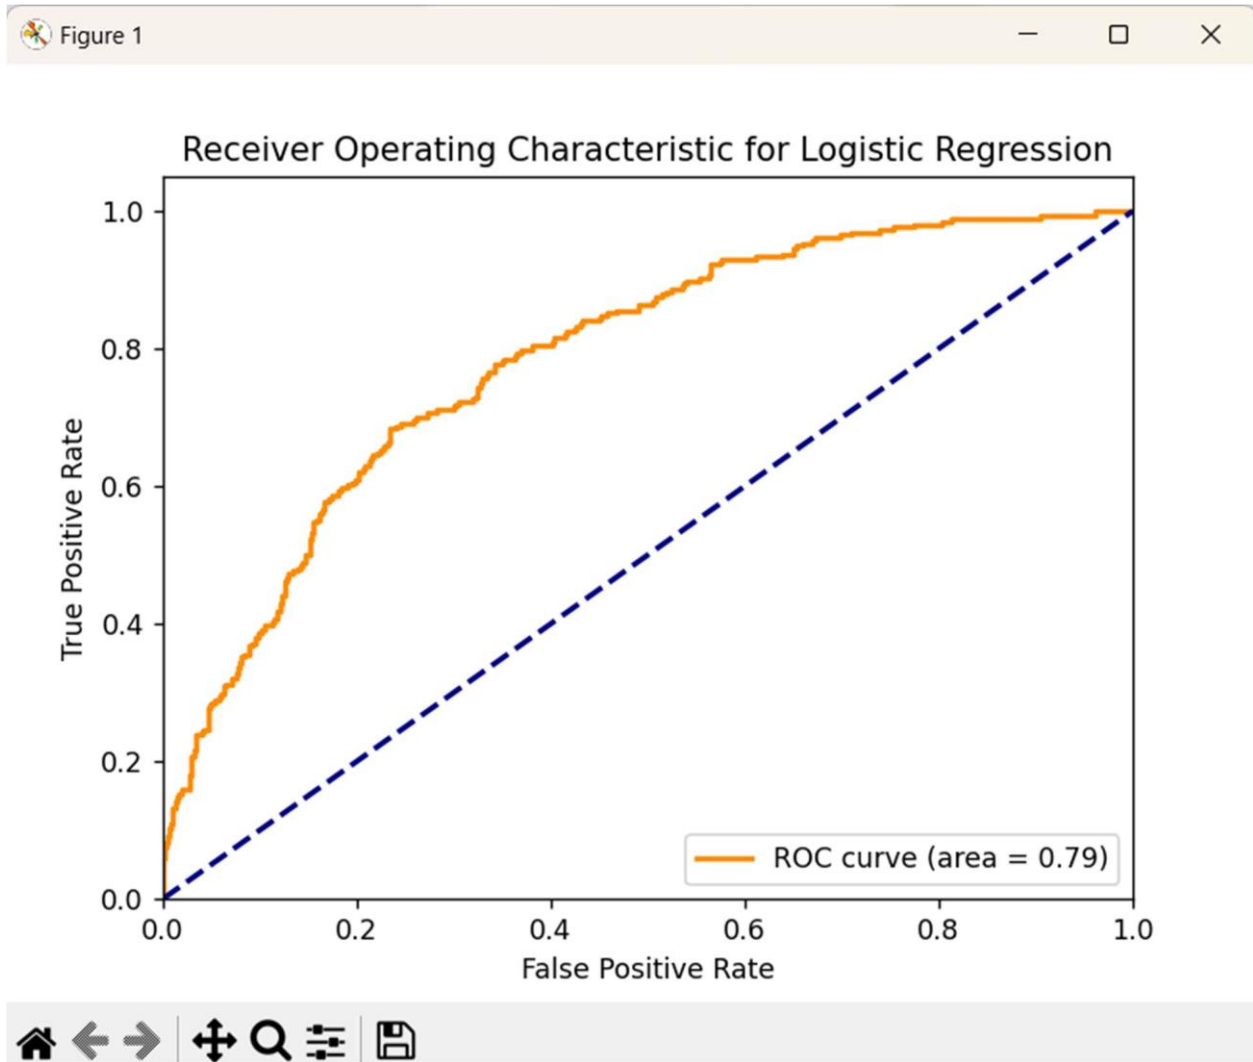
```
Logistic Regression Model Accuracy: 0.7602150537634409  
AUC Score: 0.78573581231454
```

Model Effectiveness:

- In terms of accuracy, the Logistic Regression model performed better than the Decision Tree, achieving an AUC score of 0.7857 and an approximate 76.02% superiority.

ROC Curve Analysis:

- With an AUC of 0.79, the ROC curve for the Logistic Regression model shows improved discriminative capacity. In comparison to the Decision Tree, the curve is closer to the top-left corner, indicating a greater true positive rate for most thresholds and a smaller false positive rate.



Insights and Conclusions:

- At the specified threshold of IMDb scores, the higher AUC score of the Logistic Regression model suggests that it is more successful in differentiating between 'good' and 'not good' movies. Despite the possible non-linearity in the data relationships, the enhanced discriminative performance of Logistic Regression indicates that the linear decision boundary is suitable for this dataset.
- Although the two models' accuracy scores are similar, the greater AUC of the logistic regression suggests that accuracy alone may not be the most useful statistic for assessing the model, particularly in cases where the data is unbalanced.

Linear Regression:

Justification:

- We chose linear regression, as it predicts the continuous output variables based on the independent input variables.
- In our scenario, the target variable is *imdb_score* as it ranges from 0 – 1, and we have independent variables like *release_year*, *imdb_votes*, *runtime_hours*, and different *genres*.
- We wanted to check the linear relationship between our target variable and all the independent variables to determine the best model for our dataset.

Work Done:


- Splitting the dataset into training and testing set in 80-20% ratio.
- We selected the normalized version of our dataset by removing *title*, *production_countries*, and *actor's* columns.
- Linear regression does not require any external hyperparameters to train.

Effectiveness:

- We have calculated two performance metrics, Mean Squared Error (MSE) and R-squared (R^2):
 - Mean Squared Error (MSE): It is a function that measures how close the data points are to the regression line. We obtained 0.009243, which is quite low, and it represents that our predictions are close to our actual values.
 - R-squared (R^2): It is a function that calculates how well the data fits with the regression model. We obtained 0.5405 which means it is 54.05% variance in the *imdb_score* can be explained by the model. As less than half of the model is unexplained, this may be due to non- linear relationships between some features and *imdb_score*.



- This is the heatmap of all the features correlated to each other, representing which features are strongly associated with each other.
- Higher values show they are more related to each other and can be used further for data processing. Lower values represent those features and not that co-related.



- As we can see that many data points are concentrated in the middle and for some data points, the error is too high, meaning that those points are not linearly related.
- This model effectively analyzes and provides how features are linearly related in providing higher IMDB scores, however it does not capture non-linear relationships.

Random Forest Regression:

Justification:

- Random Forest regression is an ensemble learning method which creates multiple decision trees during training and provides an output of the average of all the predictions of all trees.
- As we saw above, there are some non-linear relationships between the target column and feature columns.
- Random Forest can check non-linear relationships between the features and the target variable. This may provide a better understanding or fit for our dataset.
- Random forest reduces overfitting, as it averages multiple decision trees.
- It works for both categorical and continuous variables.

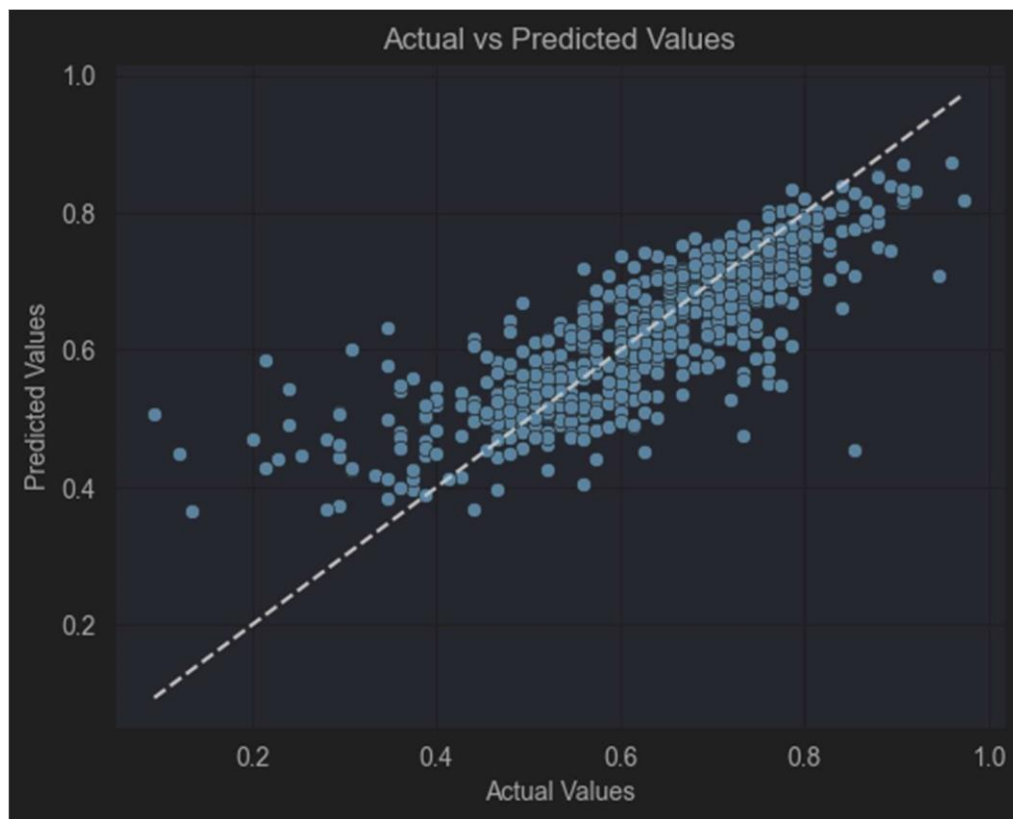
Work Done:

- Splitting the dataset into training and testing set in 80-20% ratio.
- We selected the normalized version of our dataset by removing *title*, *production_countries*, and *actor's* columns.
- Random Forest has various hyperparameters:
 - *n_estimators* (Number of Trees)
 - *max_depth* (Maximum Depth of Each Tree)
 - *min_samples_split* (Minimum Number of Samples required to Split an Internal Node)
 - *min_samples_leaf* (Minimum Number of Samples required at a leaf node)
 - *max_features* (Number of features to consider when splitting)
- To tune the hyperparameters, we used trial and error method by manually increasing and decreasing each hyperparameter, and finalizing on the hyperparameters which gave us the least Mean Square Error and High R-squared. Some of the observations are mentioned below:

<i>n_estimators</i>	<i>max_depth</i>	<i>min_samples_split</i>	<i>min_samples_leaf</i>	<i>max_features</i>	<i>MSE</i>	<i>R</i> ²
100	None	2	1	1	0.008678	56.86%
500	None	2	1	1	0.008528	57.62%
1000	None	2	1	1	0.008472	57.89%
1500	None	2	1	1	0.008451	57.99%
2500	None	2	1	1	0.008447	58.01%
1000	None	2	1	5	0.006842	65.99%
2000	17	2	1	6	0.006822	66.09%

Effectiveness:

- We have calculated two performance metrics, Mean Squared Error (MSE) and R-squared (R^2):
 - Mean Squared Error (MSE): It is a function that measures how close the data points are to the regression line. We 0.006822, which is lower than our Linear Regression model, means that it predicts the value more close than Linear Regression model.
 - R-squared (R^2): It is a function that calculates how well the data fits with the regression model. We obtained 0.66087 which means it is 66.087% variance in the *imdb_score* can be explained by the model. This higher R^2 value indicates a better fit.



- From the above graph we can see that the data points are less dispersed and are more concentrated towards this line.
- We can figure out that data points with actual values of below 0.4 are not predicted accurately. This may be due to the smaller number movies with imdb scores of 0.4 or 4 provided to our model. If we increase our dataset for such values then, this model may fit in a better way for our dataset.
- Because of its capability of handling non-linear relationships and providing average predictions of 2000 decision trees, this model is less prone to overfitting.

KMeans:

Justification:

- We chose KMeans to be able to find out more information regarding the genres of the film. To be more specific, we wanted to see if we could glean any intelligence about the combination of genres through this method.
- We believe KMeans is an effective model to find out this info as it is applied on unstructured data to gain insight into possible patterns within the data.

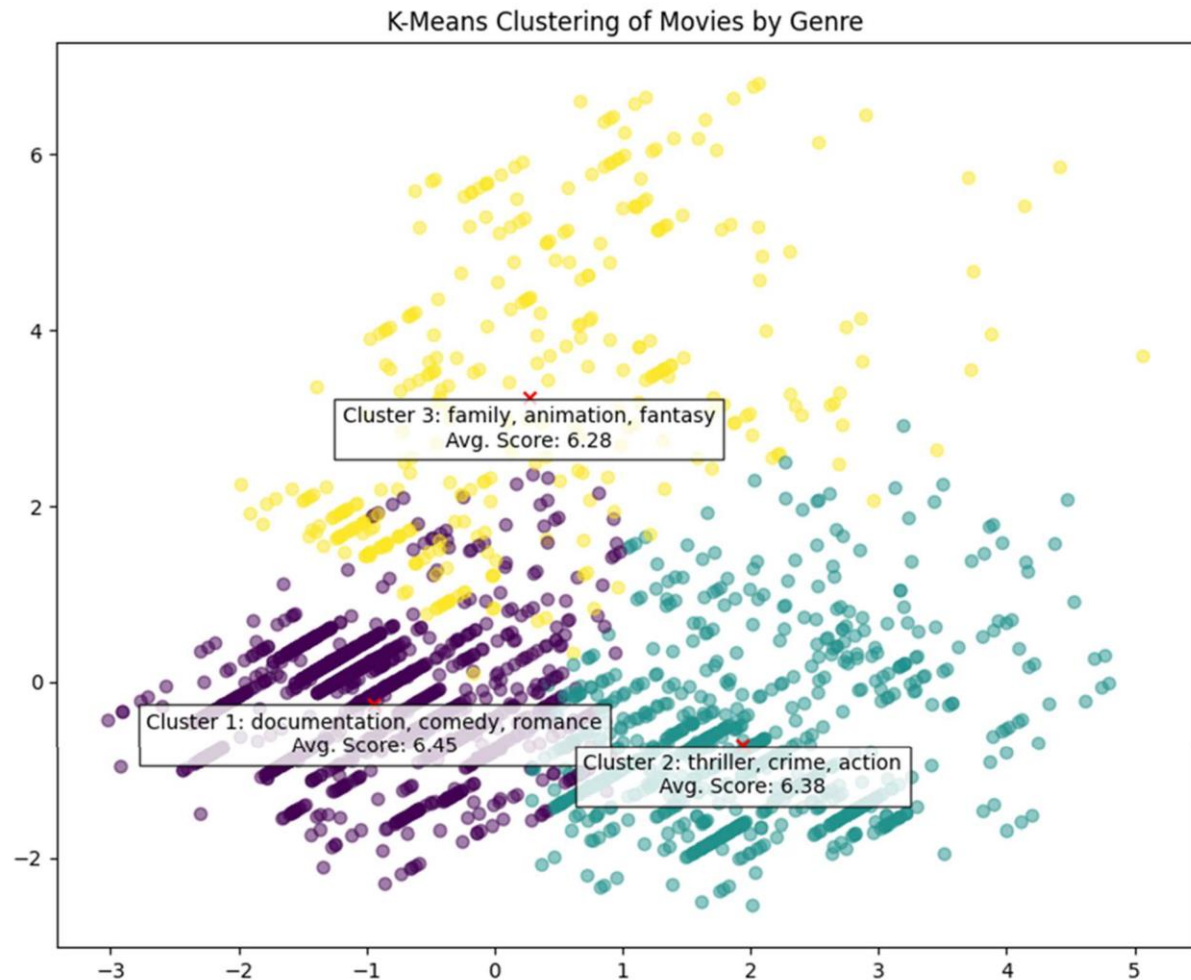
Work Done:

- As this Algorithm does not require data to be trained, we have only separated the genres from the processed dataset, standardized them, and then applied the KMeans algorithm.
- The three least common genres ('reality', 'war', 'western') have been removed.
- We then plotted the graph after reducing dimensionality to 2 through PCA. Within the clusters, we tracked the 3 most common genres and displayed those.
- Also, we have calculated the overall score of the movies in each cluster to give an idea of the relation of the score to the clustering of the genres.

Effectiveness:

- Our KMeans output visualized below reveals clusters of movies based on their genres and their average IMDb scores. We have taken $k=3$ as the best value from the elbow graph. Within the obtained 3 clusters, we see that the most common movie genres among each cluster are as: Cluster 1 – Documentation, Comedy, Romance; Cluster 2 – Thriller, Crime, Action; Cluster 3 – Family, Animation, Fantasy. Based on this info, we can see that these 3 movies are grouped together suggesting that there is prevalence of combinations of the genres in each cluster. For example, movies with a combination of genres such as Action, Crime and Action, Thriller and Thriller, Crime are observed in cluster 3. The same applies to the other 2 clusters. We also observe that average score is highest in cluster 1 which suggests that movies with genre combinations in cluster 1 such as Comedy, Romance and Documentation are the ones which are highly received hence, the higher score. Cluster 2 has the second highest score suggesting that movies with genres such as thriller, crime and action are also quite popular. From this output we can observe that the highest genres obtained per cluster are an

excellent choice of combination. Choosing such combinations would be a clever idea to obtain the highest rating.



KNN:

Justification:

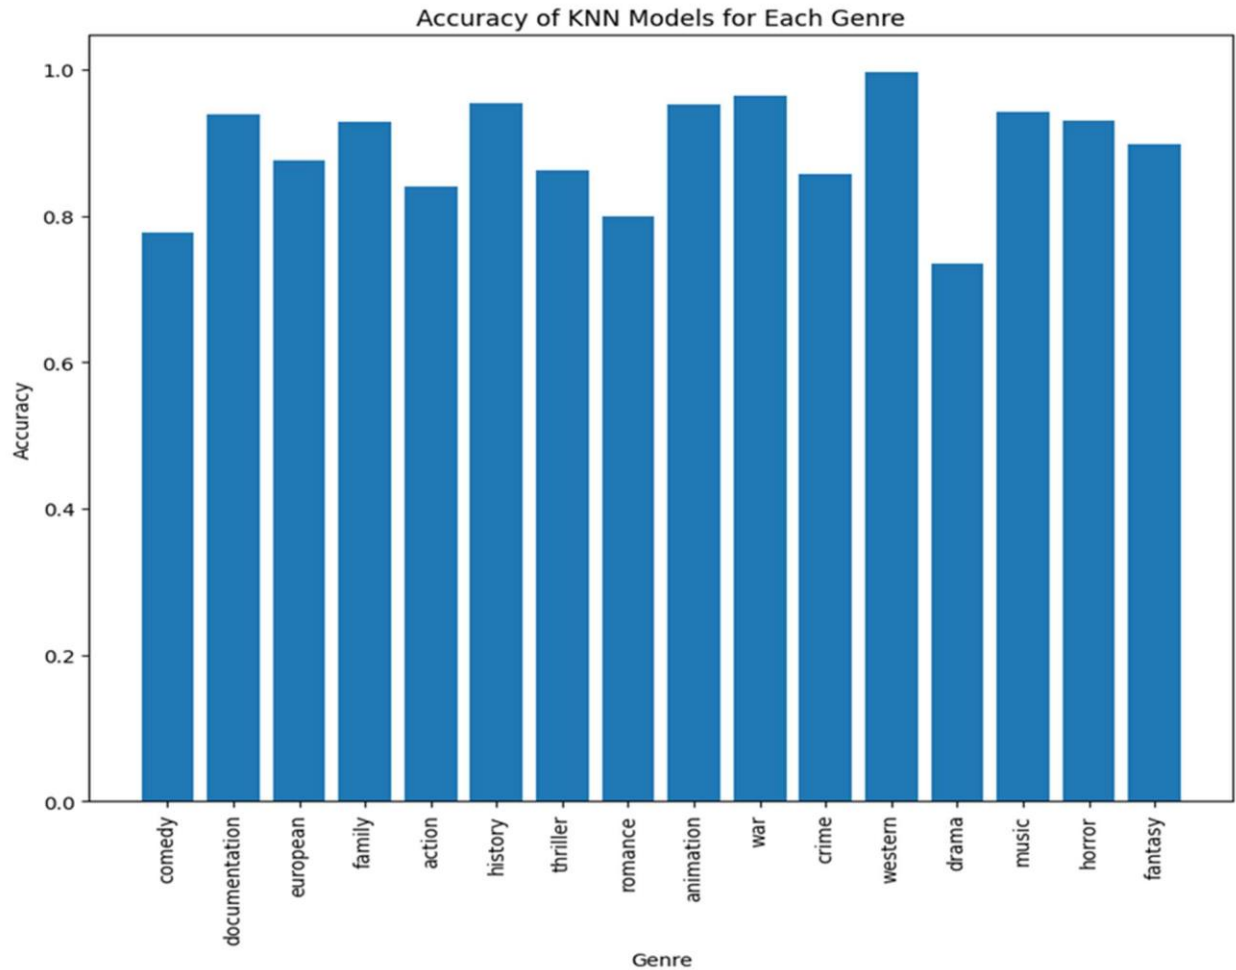
- We chose KNN here as it is a simple model which can help us to create a good identification system using its classification ability.
- Our aim here is to be able to identify the genre which in turn can be used for recommendation systems for users. Or this can also be used to correctly identify genres by production enterprises such as Netflix so that they can push forward those with greater chances of success.

Work Done:

- For this model we have taken the input to be the dataset and the target identification genre of our choice. On this basis, we have split the dataset into the target feature (genre to be identified) and the rest of the input features. After splitting the data into training and testing datasets, we trained the model using the training set and then validated the model on the testing set.
- One more thing we have done here is to try and predict each genre separately. To do this, we have provided a list of genres excluding the ones with lesser count of movies or those returning incorrect values and then iterated the model over each genre to see how well it is predicting each one.
- We have created a classification report for each genre and its separate confusion matrix as well. After plotting all of these, we have also plotted a final graph to store the accuracy obtained by the model on predicting each genre.

Effectiveness:

- We can see that our model returns predictions with quite good accuracy overall. The accuracies obtained are as below:



- Based on the above, we can say with certainty that the genre of any given movie can be predicted properly by utilizing KNN. This is immensely helpful in being able to create a recommendation system for users. It is not only in creating movies with the highest score but also in providing for the individual tastes of different users. By catering to these needs, we can properly channel the content based on the user and maximize what kinds of films can get pushed to production based on the overall popularity of each genre. Thus, this KNN application may not have been used to predict the score but by being able to successfully predict the genre, we have opened another gateway into the success of films which lies in distribution of films to users based on content.

PHASE 3

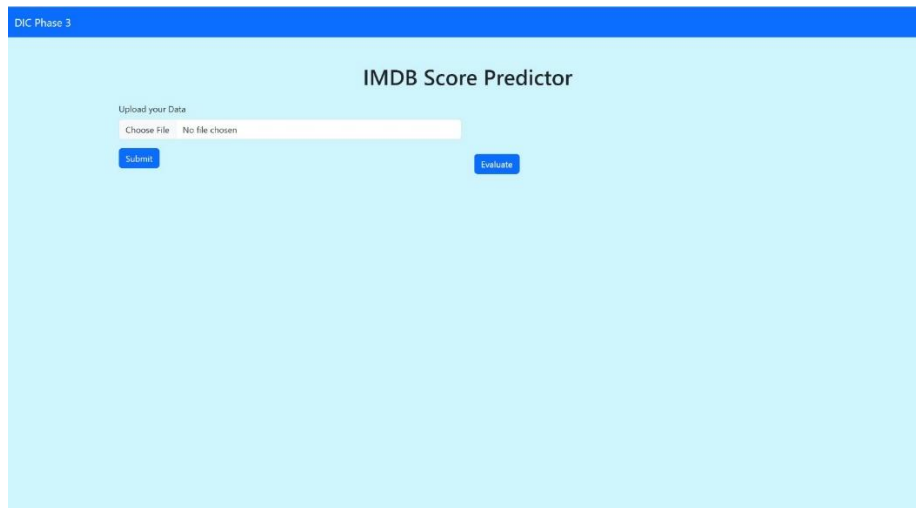
Working instructions:

1. First open code for phase-3 (recommended to open through Jupyter Notebook.)
2. Run the code.

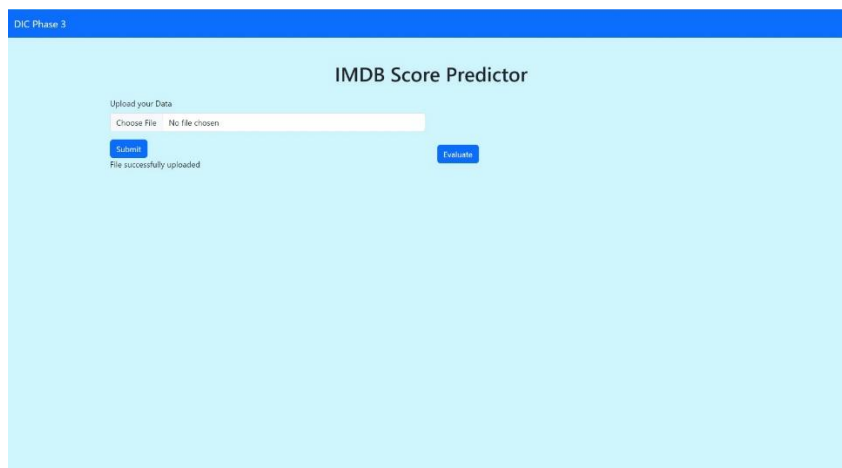
```
* Serving Flask app '__main__'
* Debug mode: off

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
```

3. Open 127.0.0.1:5000 to view the generated website.



4. Click on <Choose File> and select the dataset provided. Next click submit in order to view if file is uploaded successfully. If successful, message should appear as below:



5. In case of 'invalid file' message as shown below, please upload the corrected file in format as the provided file:

DIC Phase 3

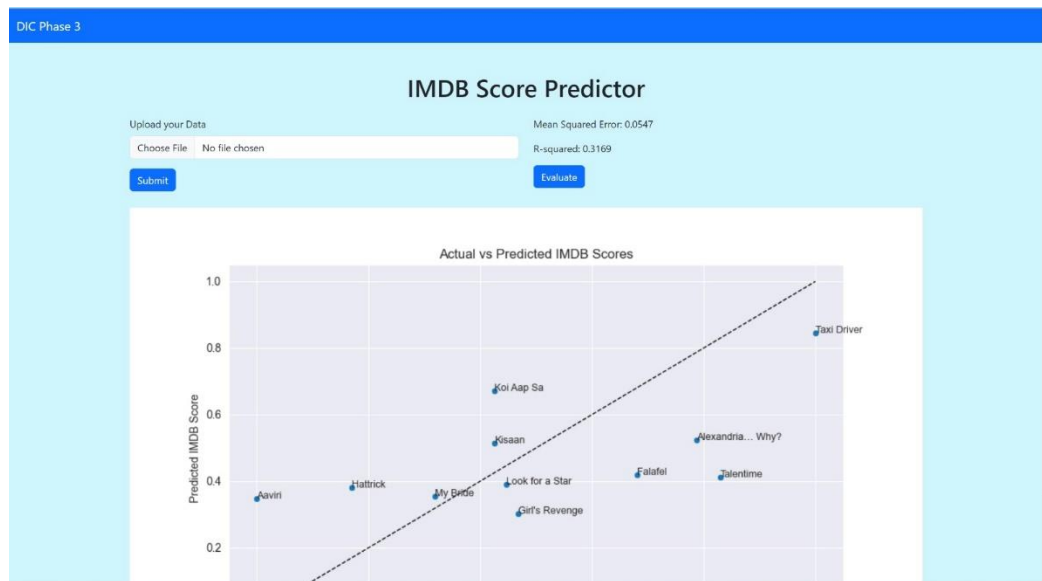
IMDB Score Predictor

Upload your Data

Choose File No file chosen

Submit Invalid File Evaluate

6. Click on evaluate and the visualization will be provided. The output showcasing the predicted evaluation metrics and visualization should appear as below:



Relevant Notes:

- In our model we have made use of the Random Forest model as that is the model which has returned the highest accuracy when predicting the IMDB score.
- A few other reasons as to why we have chosen Random Forest model is as follows:
 1. Random Forest reduces overfitting, as it averages multiple decision trees.
 2. We observed some non-linear relationships between target column and feature columns. Random forest can check non-linear relationships between the target and features so it is a well rounded choice.
 3. This model can be utilized for both categorical and continuous variables so it can work in order to predict the genre as well as the score.
- Random Forest has various hyperparameters:
 1. `n_estimators` (Number of Trees)
 2. `max_depth` (Maximum Depth of Each Tree)
 3. `min_samples_split` (Minimum Number of Samples required to Split an Internal Node)
 4. `min_samples_leaf` (Minimum Number of Samples required at a leaf node)
 5. `max_features` (Number of features to consider when splitting)
- To tune the hyperparameters, we used trial and error method by manually increasing and decreasing each hyperparameter, and finalizing on the hyperparameters which gave us the least Mean Square Error and High R-squared. Some of the observations are mentioned below:

<code>n_estimators</code>	<code>max_depth</code>	<code>min_samples_split</code>	<code>min_samples_leaf</code>	<code>max_features</code>	<i>MSE</i>	<i>R</i> ²
100	None	2	1	1	0.008678	56.86%
500	None	2	1	1	0.008528	57.62%
1000	None	2	1	1	0.008472	57.89%
1500	None	2	1	1	0.008451	57.99%
2500	None	2	1	1	0.008447	58.01%
1000	None	2	1	5	0.006842	65.99%
2000	17	2	1	6	0.006822	66.09%

- We have returned the Mean-square and *R*² error along with a visualization of the predicted value against the actual value in order to estimate the working of the model.

Recommendations related to the problem statement:

Problem Statement:

Given a streaming platform such as Netflix, we want to identify which kinds of films go on to have great ratings. Based on things such as Run time, Genre, Actors, IMDB score and number of Votes, we want to see what we can use to predict movies with a good score.

Questions we are trying to answer:

- What determines this score? The Genre, Actors, Number of people who Voted/Watched it, the Run Time, etc?
- Which of these has the most value when it comes to determining this rating?
- What kind of impact does having different combinations of Genres have?

A. Why is this a significant problem?

There is an ocean of content out in the world when it comes to watchable material. Movies and TV Shows are just a few among the many ways people consume media. These attention capturing moving pictures aren't just there to entertain us but also teach us many things. However, the real world is one in which only the popular Shows and Movies come out on top. We cannot have fantastic media to watch without investors paying the bills and investors are drawn to profit. Also, the base desire of any consumer is to have not just one amazing show to consume, but to have multiple such shows readily available. Just from these points we can see that there is an intense need to be able to filter out the best possible stories and provide those to consumers as they are the most likely to be watched. After all many people pay attention to a Movies 'score' or 'rating' before they try to watch it!

B. What is the potential of your project to contribute to the problem domain and why is it crucial?

Our project can help contribute to the problem of 'selection' and 'prediction'. Our idea is to be able to use the readily available data to be able to choose from among sample movies the ones which would be the best rated and therefore most valued. Being able to choose this would involve looking at statistics of what makes a Movie good from the dataset we have chosen and using ML to be able to train a model to be able to pick out similar such shows. This would give a clearer idea on which films can be chosen and pushed to consumers.

Benefit to Users:

What users can learn from this is of course, the output predictions of the IMDB score. They can use this completed product to be able to gauge what the IMDB score a particular movie or multiple movies may be based on the input data provided.

Users can use this functionality to find out the score of new movies which may not have a score value within IMDB although the likelihood of that is quite low.

An interesting use though, is for Users to be able to input the idea for a movie (with its relevant details filled out in the input excel sheet) and then see how this movie will do when compared to all the other movies the model was trained on. In other words, this product can help users to estimate a score for new movies or works still in their creation phase. Users can also use this model as a different source of score prediction instead of the many and often unreliable sources out there. This may help benefit users when it comes to choosing between various movies they want to watch or when they want to be able gauge the score of a movie.

A few examples of the problems that can be solved are summarized as follows:

1. User needs to choose a particular movie to watch among multiple similar movies. He can use the model to gauge the IMDB score and pick out the best.
2. User needs to filter out what new Ideas can be pushed into production. Here, user can input each idea in format of movie to the model and estimate the score.
3. User wants to pick out newly released movies in the market and choose one among them to watch.
4. User wants to re-estimate a score of very old movies using the model and see how it would compare to its previous score.

How to Extend this project:

What users can learn from this is of course, the output predictions of the IMDB score. They can use this completed product to be able to gauge what the IMDB score a particular movie or multiple movies

This project was implemented by training the model on a couple thousand movies from a dataset obtained in Kaggle. This dataset was taken from the Database of Netflix. Multiple models were then implemented with the focus of being able to predict the IMDB score of the input movies and see how well it would do. Among the models used, we even implemented a few cases wherein we tried to predict the Genre as well. Here are some ways in which we can extend this project:

- Train the model using datasets from other major streaming services, such as HBOMax, Amazon Prime, Disney+, etc along with box office statistics in order to account for more variability and improve model performance.
- To try to predict other attributes of the movie or even predict not just one but 2 or more genres the movie might be having.
- Try and see what additional genre could be mixed into an idea to improve the likelihood of the IMDB score being higher.

- Take other kinds of input such as summarized paragraphs of the movie plot and use NLP or sentiment analysis to further hone predictions.
- Implement this kind of model not only on English movies but on movies from different languages and cultures.
- Extend this model to include other forms of media such as TV Series, Anime, Books, etc.

Other Avenues that can be explored:

Although the primary aim revolved around IMDB score prediction, other forms of prediction could also be useful. A few are as follows:

- Model to generate personalized picks of movies for Users based on what they watch.
- Model which would compare popular movies across different times and see how old movies would fare in today's day and age.
- A model which could find weakpoints in a movie pitch to allow for further improvement.

REFERENCES:

1. <https://github.com/AlexTheAnalyst/PandasYouTubeSeries/blob/main/Pandas%20101%20-%20Data%20Cleaning%20in%20Pandas.ipynb>
2. <https://medium.com/analytics-vidhya/identifying-cleaning-and-replacing-outliers-titanic-dataset-20182a062893>
3. <https://www.geeksforgeeks.org/detect-and-remove-the-outliers-using-python/>
4. <https://www.adamsmith.haus/python/answers/how-to-remove-outliers-from-a-pandas-dataframe-in-python>
5. <https://www.itl.nist.gov/div898/handbook/eda/section1/eda11.html>
6. <https://www.analyticsvidhya.com/blog/2020/04/feature-scaling-machine-learning-normalization-standardization/>
7. <https://www.digitalocean.com/community/tutorials/normalize-data-in-python>
8. <https://www.kaggle.com/datasets/victorsoeiro/netflix-tv-shows-and-movies>
9. <https://scikit-learn.org/stable/modules/classes.html#regression-metrics>
10. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>
11. <https://towardsdatascience.com/random-forest-regression-5f605132d19d>
12. <https://www.investopedia.com/terms/r/r-squared.asp>
13. <https://statisticsbyjim.com/regression/interpret-r-squared-regression/>
14. <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>
15. https://scikit-learn.org/stable/model_selection.html
16. <https://scikit-learn.org/stable/modules/preprocessing.html>
17. <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>
18. <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>
19. <https://scikit-learn.org/stable/modules/tree.html>
20. <https://www.geeksforgeeks.org/decision-tree-implementation-python/#>
21. <https://towardsdatascience.com/implementing-a-decision-tree-from-scratch- f5358ff9c4bb>
22. <https://statisticsbyjim.com/regression/interpret-r-squared-regression/>
23. <https://www.ibm.com/topics/linear-regression>
24. https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html
25. <https://www.spiceworks.com/tech/artificial-intelligence/articles/what-is-linear-regression/>

Peer Evaluation Form for Final Group Work

CSE 587B

Group member 1: Vivian Vincent Dmello

Group member 2: Chaitanya Deepak Yeole

Group member 3: Vangala Eshwar Sai Prithvi Kiran

Evaluation Criteria	Group member 1	Group member 2	Group member 3
How effectively did your group mate work with you?	5	5	5
Contribution in writing the report	5	5	5
Demonstrates a cooperative and supportive attitude.	5	5	5
Contributes significantly to the success of the project.	5	5	5
TOTAL	20	20	20

Also please state the overall contribution of your teammate in percentage below, with total of all the three members accounting for 100% (33.33+33.33+33.33 ~ 100%):

Group member 1: 33.33%

Group member 2: 33.33%

Group member 3: 33.33%